

COMP4106 – Assignment 3

Brandon Schurman, 100857068

March 30, 2015

1 Classifying Independent Variables

To classify a binary vector of independent variables, we implement an optimal Bayesian Classification algorithm. In this system, we use the Bayesian Classifier to place a vector x into one of the four classes $\omega_1, \dots, \omega_4$. The true values of the four classes are randomly generated by the system in a preprocessing phase. Each class is represented as vector of $d = 10$ features, where each of the d features are randomly and independently (of the other nine features) assigned a probability value in the range of 0 to 1. The system uses these classes to randomly generate samples of 2,000 binary vectors. A binary sample vector contains d indices, where each index is assigned a 1 or instead a 0 with the probability defined in feature i of the corresponding class.

After the system obtains a sample of 2,000 binary vectors for each of the four classes, it then uses an 8-fold cross validation scheme for training and testing. Thus, 1,750 of the samples are used to train the classification algorithm, while the remaining 250 are used next to test and record the accuracy of the classification. The true values of the four classes are unknown to the classifier during training and testing. The training procedure thus uses the given samples to estimate the true probabilities of each feature, for each of the four classes. Next in the testing phase, the system classifies each sample vector x into one of the four estimated classes. It infers, using the bayesian classifier

$$\Pr(x \mid \omega_i) \leq \Pr(x \mid \omega_j)$$

pairwise for each of the four classes, which class has the highest probability of containing the vector x . More specifically, the probability of x being a vector from a class ω_j is calculated as

$$\Pr(x \mid \omega_j) = \prod_{i=1}^d \Pr(x_i \mid \omega_j)$$

After testing the classifier with 250 samples, we conclude that the classification is accurate to a degree of about 85% on average. We also notice that there are some edge cases where the sample space is significantly less accurate, to a degree of 50% or less, while other sample spaces may be classified with up to 100% accuracy. Since the system randomly generates the true classes, it is indeed possible that some class instances may have a set of features where each feature has a value of nearly 0.5, making the classification inherently difficult. Likewise, if the features in a class happen to be closer to 0.0 or 1.0, the classification is inherently much easier and thus much more accurate. Furthermore, increasing the sample size from 2,000 to 20,000 or even 200,000 was tested and did not appear to affect the accuracy of the classification in any significant way. When testing with the provided example datasets, the classification was in fact inaccurate, classifying only about 12% of vectors correctly.

2 Classifying Variables Using a Dependence Tree

To classify a binary vector of dependent variables, we must also infer a dependence tree to model the conditional probabilities of each feature d_i for each of the four classes $\omega_1, \dots, \omega_4$. Like in the independence case, the system first randomly generates the true conditional probabilities for each of the four classes in a preprocessing phase. The classes all follow the same dependence relation, which has been hard-coded into the system. The actual dependence relation between the features is shown in [Figure 2.1](#). The system generates a random sample of 2,000 binary vectors using the predefined dependence tree to train and test the classification.

Using an 8-fold cross-validation scheme for each of the four classes, the system using the first 1,750 samples to estimate the true dependence tree. To estimate the tree, the Mutual Information is calculated on each edge between every pair of features in the entire given sample space of binary vectors. The Mutual Information of the feature X relative to Y is given by

$$I(X; Y) = \sum_{x \in X} \sum_{y \in Y} p(x, y) \log \frac{p(x, y)}{p(x)p(y)}$$

Calculating the Mutual Information for every pair of features yields a set of weighted edges for a completely connected and undirected graph. The systems can then run Kruskal's algorithm on this graph to solve for the maximum spanning tree, which is an estimate of the true dependence tree. An example of an estimated dependence tree obtained by the system can be found at [Figure 2.2](#). Next, the system uses the estimated dependence tree structure to learn (i.e. estimate and store) the conditional probabilities of each feature. Once the conditional probabilities are obtained, the system can now begin testing the classification with the remaining 250 samples. To determine which class has the highest probability of producing the binary vector x , the classification algorithm uses a pairwise bayesian classifier

$$\Pr(x \mid \omega_i) \leq \Pr(x \mid \omega_j)$$

where the conditional probability that a vector x is from class ω_j is calculated as

$$\Pr(x \mid \omega_j) = \Pr(x_r) \cdot \prod_{\substack{i=1 \\ i \neq r}}^d \Pr(x_i \mid x_{p(i)})$$

where x_r is the root feature in the dependence tree, and $x_{p(i)}$ is the parent of the feature x_i in the dependence tree. Through testing, We conclude that the classification of binary dependant variables is accurate to a degree of about 50%. This is significantly less accurate than that of the independence classifier. Again, increasing the sample size does not seem to affect this accuracy. When testing with the provided example datasets, the classification was in fact inaccurate, classifying only about 12% of vectors correctly.

Figure 2.1: The true dependence tree




Figure 2.2: An approximated dependence tree

